

Nadja Rhodes

Machine Learning Engineer | Washington, DC Metro Area

[✉ narhodes1+res@gmail.com](mailto:narhodes1+res@gmail.com)

[iconix](#)

[nadjarhodes](#)

[iconix.github.io](#)

Work Experience

ASAPP

New York, NY

Jan 2019 – Jul 2024

- **Staff Machine Learning Engineer** (Nov 2021 – Jul 2024)

- Led ML engineering team of three, in XFN collaboration with a dozen research scientists and product engineers, to launch company's first "agentic" customer support bot for a Fortune 100 client, enabling automation of complex, natural language interactions.
- Led team of four developing company-wide ML training framework used by all (~20) models, implementing standardized pipelines, remote compute management, dataset versioning, model registry, and external tool integration.
- Architected multi-node distributed training on colocated servers using Kubernetes and Ray.io; collaborated on training dataset indexing service development for 10x faster dataset access (minutes to seconds).
- Proactively built an internal Slack bot as a passion project, providing thread summarization, internal asset ownership info, and trend detection to improve team communications.

- **Lead Machine Learning Engineer** (Jan 2019 – Nov 2021)

- Architected centralized routing layer between applications and ML services, migrating critical NLU components while orchestrating parallel model calls and unifying responses with entity handling and intent classification.
- Led end-to-end development of named entity recognition and entity highlighting service for 1st company voice client; owned domain-specific entity extraction from planning to rollout, delivering 11% increase in entity card hover rates.
- Implemented model observability and monitoring systems for tracking inference performance and automated alerting for model and service quality degradation.

Microsoft

Redmond, WA

Sep 2013 – Sep 2018

- **Software Engineer II** (Apr 2016 – Sep 2018)

- Pioneered and patented ML efforts in an incubation team of three, developing a task detection prototype for identifying actionable items in text.
- Rebuilt and open-sourced the [OneNote Web Clipper](#) on a modern, React-like framework; led client-side telemetry implementation. Led development of Web Clipper ratings prompt, increasing Chrome Web Store rating from 2.73 to 4.44 stars.
- Spearheaded GDPR compliance plan and implementation for exporting OneNote content.

- **Software Engineer** (Sep 2013 – Apr 2016)

- Led development and public launch of [OneNote.com/notebooks](#) entry point (4M MAU, 58K daily visits).
- Enhanced backend for OneNote welcome emails and produced data-driven insights on feature experimentation.

Education & Specialized Training

OpenAI Scholar

OpenAI

Summer 2018

San Francisco, CA

- Selected for highly competitive [scholarship](#) to study deep learning under [Dr. Natasha Jaques](#); focused on language modeling and generative text for NLP.
- Developed and open-sourced deepypebot, a generative language model that creates music commentary on Twitter. ([Code](#), [Talk](#))

B.S. Computer Science

Stanford University, School of Engineering

2009 – 2013

Stanford, CA

- Specialized in the Information track, focused on creating, processing, and understanding digital information.

Technical Skills

- **ML Technologies:** PyTorch, PyTorch Lightning, Hugging Face Transformers, spaCy, Jupyter, scikit-learn.
- **ML Domains:** NLP/NLU, Large Language Models, Text Embeddings, Neural Networks, Fine-Tuning, Conversational AI, Multi-Agent Systems, Information Extraction.
- **Infrastructure:** Kubernetes, Distributed Training (Ray), Cloud Computing (AWS), CI/CD, Airflow, MLflow, Docker, Git.
- **Backend Development:** aiohttp/asyncio servers, ML inference APIs, Microservices, gRPC/Protocol Buffers, REST.
- **Languages:** Python, TypeScript/JavaScript, SQL.

Career Transitions & Engagement

- **Recurse Center** (Nov 2024 – Feb 2025): Self-directed [programming retreat](#) in a peer-learning environment. Focused on web development, creative coding, and generative AI.
- **Family Sabbatical** (Jul – Nov 2024): Dedicated time for first-time parenting and career path reflection.
- **Presentations & Conferences:** OpenAI Scholars Demo Day (2018, [deepypebot](#)), RenderATL (2023), ACL (2019), NSBE (2018), GHC (2016), Y Combinator FFC (2014).
- **Internships:** [NTT Communication Science Labs](#) (2012), Google BOLD (2011), Stanford Research (2010).